

Sentiment Classification Of Movie Reviews

by

Dixitha Kasturi
Venkata Naga Abhiram Kuppa

Table of Contents

DATASET :	4
APPROACH.....	4
1 Reading data from csv.....	4
2 Preprocessing and filtering Data.....	5
2.1 Converting to lowercase:.....	5
2.2 Removing punctuation marks:.....	5
2.3 Removing stop words:.....	5
2.4 Filtering word tokens:.....	6
3 Generating Feature sets.....	6
3.1 Bag of Words :	7
3.2 Unigram :	7
3.3 Bigram	7
3.4 POS tagging	8
3.5 Not features	9
3.6 SL features(Sentiment lexicon)	10
3.7 LIWC features	11
3.8 Combination of LIWC and SL:.....	12
4 Saving Feature sets to CSV files :	13
Experiments :	15
5 Running Cross Validation on all Feature sets	15
5.1 Unigram:	18
5.2 Bigram:	20
5.3 POS :	22
5.4 NOT	24
5.5 SL:	26
5.6 LIWC:.....	28
5.7 SL+LIWC:	30
6 Running Naïve Bayes Classifier on Feature sets:.....	32
6.1 Unigram	32
6.2 Bigram	33
6.3 POS:	33
1.3 NOT	34
6.4 SL:	34

1.4	LIWC:.....	35
6.5	SL+LIWC:	35
7	Running Random Forest Classifier on all feature sets :	36
	Summary/Comparisons:.....	37
	Challenges	37

DATASET :

This dataset was created for the Kaggle competition, which can be found here: <https://www.kaggle.com/c/sentiment-analysis-on-movie-reviews> , and it incorporates data from Socher et al's sentiment analysis, which can be found at <http://nlp.stanford.edu/sentiment/> .

The data came from Pang and Lee's original movie review corpus, which was based on Rotten Tomatoes reviews. Socher's team used crowd-sourcing to manually mark all of the sub-sentences with a sentiment label that ranged from "negative," "little negative," "neutral," "little positive," and "positive."

The data is split into training and testing The phrases and their related sentiment labels can be found in train.tsv. The file test.tsv just contains phrases without labels. Each sentence must be given a sentiment label.

The following are the sentiment labels:

- 0 - negative
- 1 – little negative
- 2 - neutral
- 3 – little positive
- 4 – positive

APPROACH

Step 1: Reading Data from CSV file

Step 2: Preprocessing and filtering Data

Step 3: Generating Feature sets(included new features and combined features)

Step 4 : Saving Featuresets to CSV files

Step 5: Running cross validation on all Feature sets

Step 6: Running Naïve Bayes Classifier on Feature sets

Step 7: Running Random Forest Classifier on Feature sets(advanced experiment)

1 [Reading data from csv](#)

The main function takes 2 system arguments when executed. The first one is the path to the directory where the train and test files are at, the second one is a number which represents the sample size. Within this function, the function processkaggle takes these arguments, does initial processing like splitting the file into lines to further call the preprocessing function and feature set functions.

```
if __name__ == '__main__':  
    if (len(sys.argv) != 3):
```

```
print ('usage: classifyKaggle.py <corpus-dir> <limit>')
sys.exit(0)
processkaggle(sys.argv[1], sys.argv[2])

def processkaggle(dirPath,limitStr):
    # convert the limit argument from a string to an int
    limit = int(limitStr)

    os.chdir(dirPath)

    f = open('./train.tsv', 'r')
    # loop over lines in the file and use the first limit of them
    phrasedata = []
    for line in f:

        # ignore the first line starting with Phrase and read all lines
        if (not line.startswith('Phrase')):
            # remove final end of line character
            line = line.strip()
            # each line has 4 items separated by tabs
            # ignore th
            # e phrase and sentence ids, and keep the phrase and sentiment
            phrasedata.append(line.split('\t')[2:4])
```

2 Preprocessing and filtering Data

Both processed and unprocessed data was considered for all the experiments that were done.

2.1 Converting to lowercase:

Each line was converted into lowercase and split into tokens

```
w = re.split('\s+',line.lower())
```

2.2 Removing punctuation marks:

For each of the split tokens, the ones that classify as punctuations are removed from the list by substituting them with nothing.

```
punc = re.compile(r'[!#$%&()*+,-./:;<=>?@[\\]^_`{|}~]')
words = [punc.sub("",word) for word in w]
```

2.3 Removing stop words:

The predefined nltkstopwords list, some additional list of words were added that would classify as stopwords

```
words_final = []
for i in words:
```

```
if i in stopwords:
    continue
else:
    words_final.append(i)
```

2.4 Filtering word tokens:

A separate function (*filter_tokens2()*) was written to filter out tokens that were less than 2 in length, because we found a few words like 'em, 'nt etc which did not make sense.

```
def filter_token2(line):
    li=[]
    #print("line 0 ",line[0])
    for i in line[0]:
        if len(i)>2:
            li.append(i)
    return (li,line[1])
```

3 Generating Feature sets

To generate feature sets for both preprocessed and unprocessed data, different methods listed below were used. For each of the features generated, different experiments were conducted

Generating two lists of preprocessed and unprocessed tokens :

```
phrasedocs_withpre = []
phrasedocs_withoutpre= []
# add all the phrases
for phrase in phraselist:

    #Without preprocessing
    tokens = nltk.word_tokenize(phrase[0])
    phrasedocs_withoutpre.append((tokens, int(phrase[1])))

    #With preprocessing
    phrase[0] = preprocessing(phrase[0])
    tokens = nltk.word_tokenize(phrase[0])
    phrasedocs_withpre.append((tokens, int(phrase[1])))
```

Generating Filtered list for Preprocessed tokens and list for unprocessed tokens:

```
phrasedocs_withpre_filter=[]
# filtering with preprocessing and length > 2
for phrase in phrasedocs_withpre:
    phrasedocs_withpre_filter.append(filter_token2(phrase))
```

```
filtered_tokens =[]
unfiltered_tokens = []
for (d,s) in phrasedocs_withpre_filter:
    for i in d:
        filtered_tokens.append(i)

for (d,s) in phrasedocs_withoutpre:
    for i in d:
        unfiltered_tokens.append(i)
```

3.1 Bag of Words :

```
def bagOfWords(wordlist):
    wordlist = nltk.FreqDist(wordlist)
    word_feature = [w for (w,c) in wordlist.most_common(100)]
    return word_feature
```

```
filtered_bow_features = bagOfWords(filtered_tokens)
unfiltered_bow_features = bagOfWords(unfiltered_tokens)
```

3.2 Unigram :

For the document/reviews, unigram features are extracted, the features have a label "V_labelname". We take all words and turn them into Features

```
def unigram_features(doc,word_features):
    doc_words = set(doc)
    features = {}
    for word in word_features:
        features['V_%s'% word] = (word in doc_words)
    return features
```

Unigram features are extracted for both filtered and unfiltered tokens.

```
filtered_unigram_features = [(unigram_features(d,filtered_bow_features),s) for
(d,s) in phrasedocs_withpre]
unfiltered_unigram_features = [(unigram_features(d,unfiltered_bow_features),s)
for (d,s) in phrasedocs_withoutpre]
```

3.3 Bigram

The below two functions are for getting out the bigrams that are present in the data given. The function bigram_bow is like the bigram function we have learnt in class where we find the bigram collections from words in the wordlist with window size 3 and then apply filter to get bigram of only those words which have a frequency count of 6 and higher with the apply_freq_filter function. Then we get the bigrams with the chi squared to find the nbest

bigrams features which is returned. The bigram feature extraction's function i.e `bigram_features` is the snippet we have used in one of our labs to extract bigram features. This functions will be applied for both unfiltered and filtered data to compare the results.

```
def bigram_bow(wordlist):
    bigram_measure = nltk.collocations.BigramAssocMeasures()
    finder = BigramCollocationFinder.from_words(wordlist,window_size=3)
    finder.apply_freq_filter(6)
    b_features = finder.nbest(bigram_measure.chi_sq,3000)
    return b_features[:300]

def bigram_features(doc,word_features,bigram_feature):
    doc_words = set(doc)
    doc_bigrams = nltk.bigrams(doc)
    features = {}

    for word in word_features:
        features['V_{}'.format(word)] = (word in doc_words)

    for b in bigram_feature:
        features['B_{}_{}'.format(b[0],b[1])] = (b in doc_bigrams)

    return features
```

Bigram features Extraction using the above function for filtered and unfiltered data

```
filtered_bow_bigram_features=bigram_bow(filtered_tokens)
unfiltered_bow_bigram_features=bigram_bow(unfiltered_tokens)

filtered_bow_bigram_features=bigram_bow(filtered_tokens)
unfiltered_bigram_features =
[(bigram_features(d,unfiltered_bow_features,unfiltered_bow_bigram_features),s) for
(d,s) in phrasedocs_withoutpre]
```

3.4 POS tagging

Part-of-speech tagged features were extracted, Counts of various sorts of word tags are the most popular way to use POS tagging information. the function below is a feature function that counts nouns, verbs, adjectives, and adverbs.

```
def POS_features(document, word_features):
    document_words = set(document)
    tagged_words = nltk.pos_tag(document)
    features = {}
    for word in word_features:
```



```

        features['contains({})'.format(word)] = (word in document_words)
    numNoun = 0
    numVerb = 0
    numAdj = 0
    numAdverb = 0
    for (word, tag) in tagged_words:
        if tag.startswith('N'): numNoun += 1
        if tag.startswith('V'): numVerb += 1
        if tag.startswith('J'): numAdj += 1
        if tag.startswith('R'): numAdverb += 1
    features['nouns'] = numNoun
    features['verbs'] = numVerb
    features['adjectives'] = numAdj
    features['adverbs'] = numAdverb
    return features

    filtered_pos_features = [(POS_features(d,filtered_bow_features),s) for (d,s) in
    phrasedocs_withpre]

    unfiltered_pos_features = [(POS_features(d,unfiltered_bow_features),s) for (d,s)
    in phrasedocs_withoutpre]
```

3.5 Not features

Negation or contradictions in opinions should be considered when performing sentiment classification. The list of negation words considered are listed below. We can expand the list. One of the methods to analyze negation words is to find the negation word and then negate the word that follows it. This was used in the function defined below. Extracting all features and if a word is preceded by a negation word, then we negate the current word and take it as a feature.

```

negationwords = ['no', 'not', 'never', 'none', 'nowhere', 'nothing', 'noone',
'rather', 'hardly', 'scarcely', 'rarely', 'seldom', 'neither',
'nor','oppressive','revenge','revolting','rude','sticky','sick',
'sorry','tense','ugly','unwanted','unwelcome','pain','vile','vicious','quit','nons
ense','guilty','impossible','hate','damage','dead','alarming','angry','annoy','cor
rupt','creepy','cruel','cry','dishonest','dirty','evil',
'enraged','reject','sad','terrifying','stupid','yell']

def NOT_features(document, new_word_features, negationwords):
    features = {}
    for word in new_word_features:
        features['V_{}'.format(word)] = False
        features['V_NOT{}'.format(word)] = False
    # go through document words in order
    for i in range(0, len(document)):
        word = document[i]
```

```

        if ((i + 1) < len(document)) and ((word in negationwords) or
(word.endswith("\n't"))):
            i += 1
            features['V_NOT{}'.format(document[i])] = (document[i] in
new_word_features)
        else:
            features['V_{}'.format(word)] = (word in new_word_features)
    return features

```

For both filtered and unfiltered tokens, negated features were extracted.

```

filtered_not_features = [(NOT_features(d, filtered_bow_features, negationwords),
c) for (d, c) in phrasedocs_withpre]

unfiltered_not_features = [(NOT_features(d, unfiltered_bow_features,
negationwords), c) for (d, c) in phrasedocs_withoutpre]

```

3.6 SL features(Sentiment lexicon)

First, read subjective words from a subjective lexicon file prepared as part of the MPQA project by Janice Wiebe and her team at the University of Pittsburgh. These words are frequently employed as a function or in combination with other information, but they serve two purposes, one of which is to count the number of subjective positive and negative words included in each sentence document. Each word is matched to its counterpart in a list that includes both intensity and polarity. Each weak subjective word contains both positive and negative subjective words. Each strong subjective word is counted twice, and each strong subjective word is counted once. Positive count and negative count are kept track of.

```

def SL_features(document, word_features, SL):
    document_words = set(document)
    features = {}
    for word in word_features:
        features['V_{}'.format(word)] = (word in document_words)
    # count variables for the 4 classes of subjectivity
    weakPos = 0
    strongPos = 0
    weakNeg = 0
    strongNeg = 0
    for word in document_words:
        if word in SL:
            strength, posTag, isStemmed, polarity = SL[word]
            if strength == 'weaksubj' and polarity == 'positive':
                weakPos += 1
            if strength == 'strongsubj' and polarity == 'positive':
                strongPos += 1

```

```

        if strength == 'weaksubj' and polarity == 'negative':
            weakNeg += 1
        if strength == 'strongsubj' and polarity == 'negative':
            strongNeg += 1
        features['positivecount'] = weakPos + (2 * strongPos)
        features['negativecount'] = weakNeg + (2 * strongNeg)

    if 'positivecount' not in features:
        features['positivecount'] = 0
    if 'negativecount' not in features:
        features['negativecount'] = 0

    return features

```

For both filtered and unfiltered tokens these features were generated.

```

filtered_sl_features = [(SL_features(d, filtered_bow_features, SL), c) for (d,
c) in phrasedocs_withpre]

unfiltered_sl_features = [(SL_features(d, unfiltered_bow_features, SL), c) for
(d, c) in phrasedocs_withoutpre]

```

3.7 LIWC features

Linguistic Inquiry and Word Count(LIWC) is a text analysis program that calculates the percentage of words in each text that fall into one or more of over 80 linguistic, psychological and topical categories indicating various social, cognitive, and affective processes. Here we try to label the text to either a positive tweet or a negative one based on the sentiment_read_LIWC_pos_neg_words.py package given which returns words in positive emotion class and words in negative emotion class. The parameters poslist and neglist are initialized from the SL Lexicon tiff file given and brings out the positive, neutral and negative lists using which LIWC features on pos and neg words are extracted below. This function is then used to extract features for both the filtered and unfiltered data.

```

def liwc_features(doc,word_features,poslist,neglist):
    doc_words = set(doc)
    features= {}

    for word in word_features:
        features['contains({})'.format(word)] = (word in doc_words)

    pos = 0
    neg = 0
    for word in doc_words:
        if sentiment_read_LIWC_pos_neg_words.isPresent(word,poslist):
            pos+=1
        elif sentiment_read_LIWC_pos_neg_words.isPresent(word,neglist):
            neg+=1

```

```

    features ['positivecount'] = pos
    features ['negativecount'] = neg
    if 'positivecount' not in features:
        features['positivecount'] = 0
    if 'negativecount' not in features:
        features['negativecount'] = 0

    return features

```

LIWC features Extraction using the above function for filtered and unfiltered data

```

    filtered_liwc_features = [(liwc_features(d, filtered_bow_features,
poslist,neglist), c) for (d, c) in phrasedocs_withpre]

    unfiltered_liwc_features = [(liwc_features(d, unfiltered_bow_features,
poslist,neglist), c) for (d, c) in phrasedocs_withoutpre]

```

3.8 Combination of LIWC and SL:

Both LIWC AND SL features extraction methods were combined, here the strong positive and strong negative features were counted twice, by LIWC and by SL. Whereas weak positive and weak negative feature are counted only through SL feature method.

```

def combo_sl_liwc_features(doc,word_features,SL,poslist,neglist):
    doc_words = set(doc)
    features={}

    for word in word_features:
        features['contains({})'.format(word)] = (word in doc_words )

    weakPos = 0
    strongPos = 0
    weakNeg = 0
    strongNeg = 0
    for word in doc_words:
        if sentiment_read_LIWC_pos_neg_words.isPresent(word,poslist):
            strongPos +=1
        elif sentiment_read_LIWC_pos_neg_words.isPresent(word,neglist):
            strongNeg +=1
        elif word in SL:
            strength, posTag, isStemmed, polarity = SL[word]
            if strength == 'weaksubj' and polarity == 'positive':
                weakPos += 1
            if strength == 'strongsubj' and polarity == 'positive':
                strongPos += 1
            if strength == 'weaksubj' and polarity == 'negative':
                weakNeg += 1
            if strength == 'strongsubj' and polarity == 'negative':

```

```

        strongNeg += 1
    features['positivecount'] = weakPos + (2 * strongPos)
    features['negativecount'] = weakNeg + (2 * strongNeg)

    if 'positivecount' not in features:
        features['positivecount'] = 0
    if 'negativecount' not in features:
        features['negativecount'] = 0

    return features

```

For both filtered and unfiltered tokens, features were generated.

```

    filtered_combo_features = [(combo_sl_liwc_features(d, filtered_bow_features, SL,
poslist, neglist), c) for (d, c) in phrasedocs_withpre]

    unfiltered_combo_features = [(combo_sl_liwc_features(d,
unfiltered_bow_features, SL, poslist, neglist), c) for (d, c) in
phrasedocs_withoutpre]

```

4 Saving Feature sets to CSV files :

All featuresets generated so far have been saved into csv files so they can be used as training sets for any other classifier or in a different python notebook. Because of computational issues, we were not able to use these featuresets in a different python script.

A function to save the csv files was written:

```

def savingfeatures(features, path):
    f = open(path, 'w')
    featurenames = features[0][0].keys()
    fnewline = ''
    for fname in featurenames:
        fname = fname.replace(',', 'COM')
        fname = fname.replace('"', 'SQ')
        fname = fname.replace("'", 'DQ')
        fnewline += fname + ','
    fnewline += 'Level'
    f.write(fnewline)
    f.write('\n')
    for fset in features:
        featureline = ''
        for key in featurenames:
            featureline += str(fset[0][key]) + ','
        if fset[1] == 0:
            featureline += str("-1lev")
        elif fset[1] == 1:
            featureline += str("-2lev")
        elif fset[1] == 2:
            featureline += str("0lev")

```

```

elif fset[1] == 3:
    featureline += str("2lev")
elif fset[1] == 4:
    featureline += str("1lev")
f.write(featureline)
f.write('\n')
f.close()

```

To save features, the lines below were used for all features

```

savingfeatures(filtered_unigram_features, 'filtered_unigram.csv')
savingfeatures(unfiltered_unigram_features, 'unfiltered_unigram.csv')
savingfeatures(filtered_bigram_features, 'filtered_bigram.csv')
savingfeatures(unfiltered_bigram_features, 'unfiltered_bigram.csv')

```

filtered_combo.csv	5/2/2022 11:06 AM	Microsoft Excel Com...	1,036 KB
unfiltered_combo.csv	5/2/2022 11:06 AM	Microsoft Excel Com...	1,034 KB
filtered_liwc.csv	5/2/2022 11:06 AM	Microsoft Excel Com...	1,036 KB
filtered_not.csv	5/2/2022 11:06 AM	Microsoft Excel Com...	2,061 KB
filtered_sl.csv	5/2/2022 11:06 AM	Microsoft Excel Com...	1,033 KB
unfiltered_liwc.csv	5/2/2022 11:06 AM	Microsoft Excel Com...	1,034 KB
unfiltered_not.csv	5/2/2022 11:06 AM	Microsoft Excel Com...	2,059 KB
unfiltered_sl.csv	5/2/2022 11:06 AM	Microsoft Excel Com...	1,032 KB
filtered_bigram.csv	5/2/2022 11:05 AM	Microsoft Excel Com...	1,078 KB
filtered_pos.csv	5/2/2022 11:05 AM	Microsoft Excel Com...	1,038 KB
unfiltered_bigram.csv	5/2/2022 11:05 AM	Microsoft Excel Com...	1,547 KB
unfiltered_pos.csv	5/2/2022 11:05 AM	Microsoft Excel Com...	1,036 KB
unfiltered_unigram.csv	5/2/2022 11:05 AM	Microsoft Excel Com...	4,629 KB
filtered_unigram.csv	5/2/2022 11:05 AM	Microsoft Excel Com...	4,079 KB

Sample CSV file

contains(film)	contains(movie)	contains(one)	contains(rrb)	contains(lrb)	contains(ç)	contains(t)	contains(l)	contains(t)	contains(ç)	contains(r)	positivecc	negativecc	Level
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	1	-2lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	1	-1lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	1	-2lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE	0	0	-2lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	2lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	1	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	1	-2lev
TRUE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	2lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	1	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	-2lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	2lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	1	0	2lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	2	0	2lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE	FALSE	2	0	-2lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	0lev
FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	0	0	0lev

Experiments :

5 Running Cross Validation on all Feature sets

For all feature sets generated, cross validation was done using 5 folds. The average accuracy, precision, recall and f-score measures for all features were very close with *Combined SL-LIWC features* giving the highest average accuracy, precision, recall and f-score: 0.1064, 0.531, 0.494, 0.498 scores respectively for **unfiltered** data while similarly for **filtered** data we have scores in the form 0.1085, 0.543 ,0.495 ,0.499 respectively for each evaluation measure.

All the above cross validation functions are a result of running the data on the package given in crossval.py which contains the functions for cross validation and evaluation measure which we have learnt in our labs, which provide the accuracy scores for each batch of data that we run which is 31200 entries for each batch and a total of 5 batches. We then find the mean accuracy, precision, recall and F1 scores for each batch and collectively.

```
def cross_validation_PRF(num_folds, featuresets, labels):
    subset_size = int(len(featuresets)/num_folds)
    print('Each fold size:', subset_size)
    # for the number of labels - start the totals lists with zeroes
    num_labels = len(labels)
    total_precision_list = [0] * num_labels
    total_recall_list = [0] * num_labels
    total_F1_list = [0] * num_labels

    # iterate over the folds
    for i in range(num_folds):
        test_this_round = featuresets[(i*subset_size):][:subset_size]
        train_this_round = featuresets[: (i*subset_size)] +
        featuresets[((i+1)*subset_size):]
        # train using train_this_round
        classifier = nltk.NaiveBayesClassifier.train(train_this_round)
        # evaluate against test_this_round to produce the gold and predicted
        labels
        goldlist = []
        predictedlist = []
        for (features, label) in test_this_round:
            goldlist.append(label)
            predictedlist.append(classifier.classify(features))

        # computes evaluation measures for this fold and
        # returns list of measures for each label
        print('Fold', i)
        (precision_list, recall_list, F1_list) \
            = eval_measures(goldlist, predictedlist, labels)
        # take off triple string to print precision, recall and F1 for each fold
```

```
#calculating accuracy
accuracy_list= []
accuracy_this_round = nltk.classify.accuracy(classifier,test_this_round)
accuracy_list.append(accuracy_this_round)

print('\tPrecision\tRecall\t\tF1')
# print measures for each label
for i, lab in enumerate(labels):
    print(lab, '\t', "{:10.3f}".format(precision_list[i]), \
          "{:10.3f}".format(recall_list[i]), "{:10.3f}".format(F1_list[i]))

# for each label add to the sums in the total lists
for i in range(num_labels):
    # for each label, add the 3 measures to the 3 lists of totals
    total_precision_list[i] += precision_list[i]
    total_recall_list[i] += recall_list[i]
    total_F1_list[i] += F1_list[i]

# find precision, recall and F measure averaged over all rounds for all labels
# compute averages from the totals lists
precision_list = [tot/num_folds for tot in total_precision_list]
recall_list = [tot/num_folds for tot in total_recall_list]
F1_list = [tot/num_folds for tot in total_F1_list]

print('\nAverage Accuracy : ', sum(accuracy_list)/num_folds)
# the evaluation measures in a table with one row per label
print('\nAverage Precision\tRecall\t\tF1 \tPer Label')
# print measures for each label
for i, lab in enumerate(labels):
    print(lab, '\t', "{:10.3f}".format(precision_list[i]), \
          "{:10.3f}".format(recall_list[i]), "{:10.3f}".format(F1_list[i]))

# print macro average over all labels - treats each label equally
print('\nMacro Average Precision\tRecall\t\tF1 \tOver All Labels')
print('\t', "{:10.3f}".format(sum(precision_list)/num_labels), \
      "{:10.3f}".format(sum(recall_list)/num_labels), \
      "{:10.3f}".format(sum(F1_list)/num_labels))

# for micro averaging, weight the scores for each label by the number of items
# this is better for labels with imbalance
# first initialize a dictionary for label counts and then count them
label_counts = {}
for lab in labels:
    label_counts[lab] = 0
# count the labels
for (doc, lab) in featuresets:
    label_counts[lab] += 1
# make weights compared to the number of documents in featuresets
num_docs = len(featuresets)
```



```
label_weights = [(label_counts[lab] / num_docs) for lab in labels]
print('\nLabel Counts', label_counts)
#print('Label weights', label_weights)
# print macro average over all labels
print('Micro Average Precision\tRecall\t\tF1 \tOver All Labels')
precision = sum([a * b for a,b in zip(precision_list, label_weights)])
recall = sum([a * b for a,b in zip(recall_list, label_weights)])
F1 = sum([a * b for a,b in zip(F1_list, label_weights)])
print( '\t', "{:10.3f}".format(precision), \
      "{:10.3f}".format(recall), "{:10.3f}".format(F1))
# Function to compute precision, recall and F1 for each label
# and for any number of labels
# Input: list of gold labels, list of predicted labels (in same order)
# Output: returns lists of precision, recall and F1 for each label
#       (for computing averages across folds and labels)
def eval_measures(gold, predicted, labels):

    # these lists have values for each label
    recall_list = []
    precision_list = []
    F1_list = []

    for lab in labels:
        # for each label, compare gold and predicted lists and compute values
        TP = FP = FN = TN = 0
        for i, val in enumerate(gold):
            if val == lab and predicted[i] == lab: TP += 1
            if val == lab and predicted[i] != lab: FN += 1
            if val != lab and predicted[i] == lab: FP += 1
            if val != lab and predicted[i] != lab: TN += 1
        # use these to compute recall, precision, F1
        # for small numbers, guard against dividing by zero in computing measures
        if (TP == 0) or (FP == 0) or (FN == 0):
            recall_list.append(0)
            precision_list.append(0)
            F1_list.append(0)
        else:
            recall = TP / (TP + FP)
            precision = TP / (TP + FN)
            recall_list.append(recall)
            precision_list.append(precision)
            F1_list.append( 2 * (recall * precision) / (recall + precision))

    # the evaluation measures in a table with one row per label
    return (precision_list, recall_list, F1_list)
```

5.1 Unigram:

For unfiltered tokens:

```
print("\n Unigram Unfiltered : ")
crossval.cross_validation_PRF(5,unfiltered_unigram_features,labels)
```

```
Unigram Unfiltered :
Each fold size: 31200
Fold 0
Precision      Recall      F1
0      0.207      0.170      0.187
1      0.182      0.338      0.237
2      0.840      0.594      0.696
3      0.160      0.326      0.214
4      0.126      0.277      0.173
Fold 1
Precision      Recall      F1
0      0.211      0.161      0.182
1      0.171      0.336      0.226
2      0.847      0.598      0.701
3      0.173      0.352      0.232
4      0.119      0.262      0.163
Fold 2
Precision      Recall      F1
0      0.187      0.150      0.167
1      0.166      0.329      0.221
2      0.846      0.602      0.703
3      0.163      0.336      0.220
4      0.121      0.227      0.158
Fold 3
Precision      Recall      F1
0      0.201      0.173      0.186
1      0.174      0.336      0.229
2      0.850      0.593      0.698
3      0.159      0.329      0.214
4      0.133      0.275      0.179
Fold 4
Precision      Recall      F1
0      0.210      0.178      0.193
1      0.187      0.346      0.243
2      0.844      0.600      0.701
3      0.168      0.344      0.225
4      0.122      0.234      0.161

Average Accuracy : 0.10319871794871796

Average Precision      Recall      F1      Per Label
0      0.203      0.167      0.183
1      0.176      0.337      0.231
2      0.845      0.597      0.700
3      0.165      0.337      0.221
4      0.124      0.255      0.167

Macro Average Precision Recall      F1      Over All Labels
0.303      0.339      0.300

Label Counts {0: 7067, 1: 27264, 2: 79551, 3: 32914, 4: 9204}
Micro Average Precision Recall      F1      Over All Labels
0.513      0.457      0.462
```

For Unigram filtered tokens:

```
print("\n Unigram filtered : ")
crossval.cross_validation_PRF(5,filtered_unigram_features,labels)
```

```
Each fold size: 31200
Fold 0
      Precision      Recall      F1
0      0.051      0.346      0.089
1      0.089      0.328      0.140
2      0.927      0.559      0.697
3      0.158      0.380      0.224
4      0.065      0.359      0.111
Fold 1
      Precision      Recall      F1
0      0.050      0.320      0.086
1      0.090      0.345      0.142
2      0.931      0.558      0.698
3      0.168      0.417      0.240
4      0.067      0.320      0.111
Fold 2
      Precision      Recall      F1
0      0.048      0.322      0.083
1      0.090      0.353      0.143
2      0.928      0.557      0.696
3      0.162      0.397      0.231
4      0.066      0.340      0.110
Fold 3
      Precision      Recall      F1
0      0.055      0.330      0.095
1      0.082      0.349      0.132
2      0.929      0.555      0.695
3      0.154      0.388      0.220
4      0.062      0.302      0.104
Fold 4
      Precision      Recall      F1
0      0.053      0.294      0.090
1      0.083      0.340      0.133
2      0.928      0.554      0.693
3      0.157      0.380      0.223
4      0.071      0.405      0.121

Average Accuracy : 0.10510897435897434

Average Precision      Recall      F1      Per Label
0      0.051      0.322      0.089
1      0.087      0.343      0.138
2      0.928      0.556      0.696
3      0.160      0.393      0.227
4      0.066      0.345      0.111

Macro Average Precision Recall      F1      Over All Labels
0.259      0.392      0.252

Label Counts {0: 7069, 1: 27263, 2: 79557, 3: 32912, 4: 9199}
Micro Average Precision Recall      F1      Over All Labels
0.529      0.461      0.437
```

5.2 Bigram:

For unfiltered :

```
print("\n Bigram Unfiltered : ")
crossval.cross_validation_PRF(5,unfiltered_bigram_features,labels)
```

```
Bigram Unfiltered :
Each fold size: 31200
Fold 0
      Precision      Recall      F1
0      0.207      0.170      0.187
1      0.182      0.338      0.237
2      0.840      0.594      0.696
3      0.160      0.326      0.214
4      0.126      0.277      0.173
Fold 1
      Precision      Recall      F1
0      0.211      0.161      0.182
1      0.171      0.336      0.226
2      0.847      0.598      0.701
3      0.173      0.352      0.232
4      0.119      0.262      0.163
Fold 2
      Precision      Recall      F1
0      0.187      0.150      0.167
1      0.166      0.329      0.221
2      0.846      0.602      0.703
3      0.163      0.336      0.220
4      0.121      0.227      0.157
Fold 3
      Precision      Recall      F1
0      0.200      0.173      0.185
1      0.174      0.336      0.229
2      0.850      0.593      0.698
3      0.159      0.329      0.214
4      0.133      0.275      0.179
Fold 4
      Precision      Recall      F1
0      0.210      0.178      0.193
1      0.187      0.346      0.243
2      0.844      0.600      0.701
3      0.168      0.344      0.226
4      0.122      0.234      0.161
Average Accuracy : 0.10320512820512821

Average Precision      Recall      F1      Per Label
0      0.203      0.167      0.183
1      0.176      0.337      0.231
2      0.845      0.597      0.700
3      0.165      0.337      0.221
4      0.124      0.255      0.167

Macro Average Precision Recall      F1      Over All Labels
0.303      0.339      0.300

Label Counts {0: 7067, 1: 27264, 2: 79551, 3: 32914, 4: 9204}
Micro Average Precision Recall      F1      Over All Labels
0.513      0.457      0.462
```

For Bigram filtered tokens

```
print("\n Bigram filtered : ")
crossval.cross_validation_PRF(5,filtered_bigram_features,labels)
```

```
Each fold size: 31200
Fold 0
      Precision      Recall      F1
0      0.051      0.346      0.089
1      0.089      0.328      0.140
2      0.927      0.559      0.697
3      0.158      0.380      0.224
4      0.065      0.359      0.111
Fold 1
      Precision      Recall      F1
0      0.050      0.320      0.086
1      0.090      0.345      0.142
2      0.931      0.558      0.698
3      0.168      0.417      0.240
4      0.067      0.320      0.111
Fold 2
      Precision      Recall      F1
0      0.048      0.322      0.083
1      0.090      0.353      0.143
2      0.928      0.557      0.696
3      0.163      0.397      0.231
4      0.066      0.340      0.110
Fold 3
      Precision      Recall      F1
0      0.055      0.330      0.095
1      0.082      0.349      0.132
2      0.929      0.555      0.695
3      0.154      0.388      0.220
4      0.062      0.302      0.104
Fold 4
      Precision      Recall      F1
0      0.053      0.294      0.090
1      0.083      0.340      0.133
2      0.928      0.554      0.693
3      0.157      0.380      0.223
4      0.071      0.405      0.121

Average Accuracy : 0.10510897435897434

Average Precision      Recall      F1      Per Label
0      0.051      0.322      0.089
1      0.087      0.343      0.138
2      0.928      0.556      0.696
3      0.160      0.393      0.227
4      0.066      0.345      0.111

Macro Average Precision Recall      F1      Over All Labels
0.259      0.392      0.252

Label Counts {0: 7069, 1: 27263, 2: 79557, 3: 32912, 4: 9199}
Micro Average Precision Recall      F1      Over All Labels
0.529      0.461      0.438
```

5.3 POS :

For unfiltered :

```
print("\n Pos Unfiltered : ")
crossval.cross_validation_PRF(5,unfiltered_pos_features,labels)
```

```
Pos Unfiltered :
Each fold size: 31200
Fold 0
      Precision      Recall      F1
0      0.274      0.141      0.186
1      0.189      0.325      0.239
2      0.809      0.613      0.698
3      0.153      0.321      0.207
4      0.151      0.245      0.187
Fold 1
      Precision      Recall      F1
0      0.284      0.139      0.187
1      0.184      0.328      0.236
2      0.815      0.616      0.702
3      0.164      0.339      0.221
4      0.136      0.231      0.172
Fold 2
      Precision      Recall      F1
0      0.252      0.134      0.175
1      0.186      0.326      0.237
2      0.816      0.623      0.706
3      0.155      0.327      0.210
4      0.150      0.205      0.173
Fold 3
      Precision      Recall      F1
0      0.272      0.145      0.189
1      0.190      0.334      0.242
2      0.818      0.612      0.700
3      0.153      0.329      0.209
4      0.161      0.246      0.195
Fold 4
      Precision      Recall      F1
0      0.288      0.154      0.200
1      0.198      0.334      0.249
2      0.809      0.619      0.702
3      0.164      0.334      0.220
4      0.152      0.215      0.178

Average Accuracy : 0.10088461538461539

Average Precision      Recall      F1      Per Label
0      0.274      0.143      0.188
1      0.190      0.329      0.241
2      0.813      0.617      0.701
3      0.158      0.330      0.213
4      0.150      0.228      0.181

Macro Average Precision Recall      F1      Over All Labels
0.317      0.329      0.305

Label Counts {0: 7067, 1: 27264, 2: 79551, 3: 32914, 4: 9204}
Micro Average Precision Recall      F1      Over All Labels
0.502      0.462      0.464
```

For filtered tokens :

```
print("\n Pos Unfiltered : ")
crossval.cross_validation_PRF(5,filtered_pos_features,labels)
```

```
Pos Unfiltered :
Each fold size: 31200
Fold 0
Precision      Recall      F1
0      0.135      0.204      0.163
1      0.167      0.307      0.216
2      0.851      0.605      0.707
3      0.199      0.340      0.251
4      0.138      0.248      0.178
Fold 1
Precision      Recall      F1
0      0.141      0.195      0.164
1      0.164      0.326      0.218
2      0.856      0.604      0.708
3      0.214      0.366      0.270
4      0.133      0.224      0.167
Fold 2
Precision      Recall      F1
0      0.144      0.215      0.172
1      0.176      0.322      0.227
2      0.846      0.600      0.702
3      0.198      0.353      0.254
4      0.145      0.236      0.180
Fold 3
Precision      Recall      F1
0      0.139      0.210      0.167
1      0.160      0.315      0.212
2      0.847      0.597      0.700
3      0.197      0.341      0.250
4      0.144      0.248      0.182
Fold 4
Precision      Recall      F1
0      0.131      0.190      0.155
1      0.159      0.309      0.210
2      0.849      0.596      0.701
3      0.198      0.345      0.252
4      0.143      0.252      0.183

Average Accuracy : 0.10308333333333333

Average Precision      Recall      F1      Per Label
0      0.138      0.203      0.164
1      0.165      0.316      0.217
2      0.850      0.601      0.704
3      0.201      0.349      0.255
4      0.141      0.242      0.178

Macro Average Precision      Recall      F1      Over All Labels
0.299      0.342      0.304

Label Counts {0: 7069, 1: 27263, 2: 79557, 3: 32912, 4: 9199}
Micro Average Precision      Recall      F1      Over All Labels
0.519      0.459      0.469
```

5.4 NOT

NOT Unfiltered Tokens:

```
print("\n NOT unfiltered : ")
crossval.cross_validation_PRF(5, unfiltered_not_features, labels)
```

```
NOT filtered :
Each fold size: 31200
Fold 0
Precision      Recall      F1
0      0.595      0.228      0.329
1      0.426      0.445      0.436
2      0.642      0.767      0.699
3      0.441      0.489      0.464
4      0.562      0.333      0.418
Fold 1
Precision      Recall      F1
0      0.597      0.244      0.346
1      0.424      0.447      0.436
2      0.649      0.758      0.699
3      0.427      0.496      0.459
4      0.572      0.330      0.419
Fold 2
Precision      Recall      F1
0      0.601      0.247      0.350
1      0.428      0.445      0.436
2      0.646      0.771      0.703
3      0.444      0.492      0.467
4      0.555      0.316      0.402
Fold 3
Precision      Recall      F1
0      0.611      0.232      0.336
1      0.423      0.453      0.438
2      0.629      0.752      0.685
3      0.432      0.483      0.456
4      0.582      0.324      0.417
Fold 4
Precision      Recall      F1
0      0.594      0.237      0.339
1      0.418      0.442      0.430
2      0.636      0.752      0.689
3      0.445      0.488      0.465
4      0.561      0.348      0.429

Average Accuracy : 0.11024358974358975
```

```
Average Accuracy : 0.11024358974358975

Average Precision      Recall      F1      Per Label
0      0.600      0.237      0.340
1      0.424      0.447      0.435
2      0.641      0.760      0.695
3      0.438      0.490      0.462
4      0.566      0.330      0.417

Macro Average Precision      Recall      F1      Over All Labels
0.534      0.453      0.470

Label Counts {0: 7069, 1: 27263, 2: 79554, 3: 32914, 4: 9200}
Micro Average Precision      Recall      F1      Over All Labels
0.554      0.599      0.568
```


NOT Filtered tokens:

```
print("\n NOT filtered : ")
crossval.cross_validation_PRF(5, filtered_not_features, labels)
```

```
NOT filtered :
Each fold size: 31200
Fold 0
      Precision      Recall      F1
0      0.595      0.228      0.329
1      0.426      0.445      0.436
2      0.642      0.767      0.699
3      0.441      0.489      0.464
4      0.562      0.333      0.418
Fold 1
      Precision      Recall      F1
0      0.597      0.244      0.346
1      0.424      0.447      0.436
2      0.649      0.758      0.699
3      0.427      0.496      0.459
4      0.572      0.330      0.419
Fold 2
      Precision      Recall      F1
0      0.601      0.247      0.350
1      0.428      0.445      0.436
2      0.646      0.771      0.703
3      0.444      0.492      0.467
4      0.555      0.316      0.402
Fold 3
      Precision      Recall      F1
0      0.611      0.232      0.336
1      0.423      0.453      0.438
2      0.629      0.752      0.685
3      0.432      0.483      0.456
4      0.582      0.324      0.417
Fold 4
      Precision      Recall      F1
0      0.594      0.237      0.339
1      0.418      0.442      0.430
2      0.636      0.752      0.689
3      0.445      0.488      0.465
4      0.561      0.348      0.429

Average Accuracy : 0.11024358974358975

Average Accuracy : 0.11024358974358975

Average Precision      Recall      F1      Per Label
0      0.600      0.237      0.340
1      0.424      0.447      0.435
2      0.641      0.760      0.695
3      0.438      0.490      0.462
4      0.566      0.330      0.417

Macro Average Precision Recall      F1      Over All Labels
0.534      0.453      0.470

Label Counts {0: 7069, 1: 27263, 2: 79554, 3: 32914, 4: 9200}
Micro Average Precision Recall      F1      Over All Labels
0.554      0.599      0.568
```

5.5 SL:

For unfiltered:

```
print("\n SL Unfiltered : ")
crossval.cross_validation_PRF(5,unfiltered_sl_features,labels)
```

```
SL Unfiltered :
Each fold size: 31200
Fold 0
Precision      Recall      F1
0      0.249      0.192      0.217
1      0.210      0.368      0.267
2      0.822      0.627      0.712
3      0.247      0.398      0.305
4      0.217      0.307      0.254
Fold 1
Precision      Recall      F1
0      0.260      0.190      0.220
1      0.196      0.364      0.255
2      0.820      0.625      0.710
3      0.254      0.393      0.309
4      0.213      0.319      0.255
Fold 2
Precision      Recall      F1
0      0.233      0.177      0.202
1      0.193      0.355      0.250
2      0.824      0.632      0.715
3      0.252      0.399      0.309
4      0.213      0.285      0.244
Fold 3
Precision      Recall      F1
0      0.257      0.209      0.230
1      0.197      0.371      0.258
2      0.829      0.622      0.711
3      0.246      0.397      0.304
4      0.226      0.300      0.258
Fold 4
Precision      Recall      F1
0      0.249      0.197      0.220
1      0.212      0.373      0.271
2      0.821      0.629      0.712
3      0.258      0.402      0.314
4      0.215      0.297      0.250

Average Accuracy : 0.10691666666666666

Average Precision      Recall      F1      Per Label
0      0.250      0.193      0.218
1      0.202      0.366      0.260
2      0.823      0.627      0.712
3      0.251      0.398      0.308
4      0.217      0.301      0.252

Macro Average Precision      Recall      F1      Over All Labels
0.349      0.377      0.350

Label Counts {0: 7067, 1: 27264, 2: 79551, 3: 32914, 4: 9204}
Micro Average Precision      Recall      F1      Over All Labels
0.532      0.494      0.498
```

For filtered tokens :

```
print("\n SL Unfiltered : ")
crossval.cross_validation_PRF(5,filtered_sl_features,labels)
```

```
SL Unfiltered :
Each fold size: 31200
Fold 0
      Precision      Recall      F1
0      0.106      0.321      0.159
1      0.148      0.344      0.207
2      0.849      0.612      0.711
3      0.322      0.386      0.351
4      0.149      0.349      0.209
Fold 1
      Precision      Recall      F1
0      0.103      0.305      0.154
1      0.144      0.354      0.205
2      0.858      0.616      0.717
3      0.339      0.403      0.368
4      0.140      0.317      0.194
Fold 2
      Precision      Recall      F1
0      0.100      0.305      0.151
1      0.157      0.357      0.218
2      0.844      0.610      0.708
3      0.325      0.394      0.356
4      0.147      0.326      0.202
Fold 3
      Precision      Recall      F1
0      0.108      0.312      0.161
1      0.148      0.353      0.209
2      0.851      0.608      0.709
3      0.320      0.395      0.354
4      0.143      0.326      0.199
Fold 4
      Precision      Recall      F1
0      0.123      0.329      0.179
1      0.151      0.356      0.212
2      0.852      0.609      0.710
3      0.329      0.402      0.362
4      0.164      0.383      0.230

Average Accuracy : 0.10878205128205128

Average Precision      Recall      F1      Per Label
0      0.108      0.314      0.161
1      0.150      0.353      0.210
2      0.851      0.611      0.711
3      0.327      0.396      0.358
4      0.149      0.340      0.207

Macro Average Precision Recall      F1      Over All Labels
0.317      0.403      0.329

Label Counts {0: 7069, 1: 27263, 2: 79557, 3: 32912, 4: 9199}
Micro Average Precision Recall      F1      Over All Labels
0.543      0.491      0.494
```

5.6 LIWC:

For unfiltered :

```
print("\n LIWC Unfiltered : ")
crossval.cross_validation_PRF(5,unfiltered_liwc_features,labels)
```

```
LIWC Unfiltered :
Each fold size: 31200
Fold 0
      Precision      Recall      F1
0      0.239      0.193      0.213
1      0.199      0.360      0.256
2      0.830      0.620      0.710
3      0.239      0.398      0.298
4      0.193      0.300      0.235
Fold 1
      Precision      Recall      F1
0      0.249      0.190      0.215
1      0.188      0.349      0.244
2      0.830      0.619      0.709
3      0.238      0.396      0.297
4      0.185      0.299      0.229
Fold 2
      Precision      Recall      F1
0      0.219      0.175      0.194
1      0.189      0.350      0.245
2      0.832      0.625      0.714
3      0.230      0.388      0.289
4      0.189      0.267      0.221
Fold 3
      Precision      Recall      F1
0      0.233      0.195      0.212
1      0.187      0.355      0.245
2      0.836      0.613      0.707
3      0.228      0.392      0.288
4      0.195      0.289      0.233
Fold 4
      Precision      Recall      F1
0      0.241      0.199      0.218
1      0.199      0.360      0.257
2      0.832      0.623      0.712
3      0.237      0.395      0.296
4      0.191      0.279      0.227

Average Accuracy : 0.10632692307692308

Average Precision      Recall      F1      . Per Label
0      0.236      0.190      0.211
1      0.192      0.355      0.250
2      0.832      0.620      0.710
3      0.234      0.394      0.294
4      0.191      0.287      0.229

Macro Average Precision Recall      F1      Over All Labels
0.337      0.369      0.339

Label Counts {0: 7067, 1: 27264, 2: 79551, 3: 32914, 4: 9204}
Micro Average Precision Recall      F1      Over All Labels
0.529      0.487      0.491
```

For filtered tokens :

```
print("\n LIWC Unfiltered : ")
crossval.cross_validation_PRF(5,filtered_liwc_features,labels)
```

```
LIWC Unfiltered :
Each fold size: 31200
Fold 0
      Precision      Recall      F1
0      0.097      0.347      0.151
1      0.117      0.361      0.177
2      0.878      0.595      0.709
3      0.283      0.384      0.326
4      0.124      0.336      0.181
Fold 1
      Precision      Recall      F1
0      0.098      0.328      0.151
1      0.105      0.373      0.164
2      0.888      0.594      0.712
3      0.288      0.399      0.334
4      0.127      0.311      0.180
Fold 2
      Precision      Recall      F1
0      0.093      0.311      0.143
1      0.110      0.357      0.169
2      0.872      0.590      0.704
3      0.284      0.389      0.328
4      0.133      0.344      0.192
Fold 3
      Precision      Recall      F1
0      0.105      0.336      0.161
1      0.104      0.368      0.162
2      0.882      0.588      0.706
3      0.267      0.383      0.314
4      0.129      0.315      0.183
Fold 4
      Precision      Recall      F1
0      0.116      0.319      0.171
1      0.109      0.371      0.168
2      0.879      0.590      0.706
3      0.281      0.391      0.327
4      0.142      0.371      0.206

Average Accuracy : 0.10776282051282052

Average Precision      Recall      F1      Per Label
0      0.102      0.328      0.155
1      0.109      0.366      0.168
2      0.880      0.591      0.707
3      0.281      0.389      0.326
4      0.131      0.336      0.188

Macro Average Precision Recall      F1      Over All Labels
0.300      0.402      0.309

Label Counts {0: 7069, 1: 27263, 2: 79557, 3: 32912, 4: 9199}
Micro Average Precision Recall      F1      Over All Labels
0.539      0.482      0.477
```

5.7 SL+LIWC:

For unfiltered:

```
print("\n Combined SL LIWC Unfiltered : ")
crossval.cross_validation_PRF(5,unfiltered_combo_features,labels)
```

```
Combined SL LIWC Unfiltered :
Each fold size: 31200
Fold 0
Precision      Recall      F1
0      0.253      0.191      0.218
1      0.210      0.367      0.267
2      0.818      0.629      0.711
3      0.254      0.394      0.309
4      0.214      0.309      0.253
Fold 1
Precision      Recall      F1
0      0.263      0.186      0.218
1      0.197      0.364      0.256
2      0.818      0.626      0.710
3      0.254      0.388      0.307
4      0.200      0.308      0.242
Fold 2
Precision      Recall      F1
0      0.241      0.183      0.208
1      0.201      0.364      0.259
2      0.821      0.634      0.715
3      0.256      0.397      0.312
4      0.210      0.289      0.243
Fold 3
Precision      Recall      F1
0      0.258      0.205      0.228
1      0.195      0.366      0.254
2      0.826      0.623      0.710
3      0.247      0.390      0.303
4      0.214      0.291      0.247
Fold 4
Precision      Recall      F1
0      0.253      0.197      0.222
1      0.214      0.371      0.271
2      0.815      0.630      0.711
3      0.262      0.394      0.314
4      0.201      0.287      0.236

Average Accuracy : 0.10642307692307693

Average Precision      Recall      F1      Per Label
0      0.254      0.192      0.219
1      0.203      0.367      0.261
2      0.820      0.628      0.711
3      0.255      0.393      0.309
4      0.208      0.297      0.244

Macro Average Precision      Recall      F1      Over All Labels
0.348      0.375      0.349

Label Counts {0: 7067, 1: 27264, 2: 79551, 3: 32914, 4: 9204}
Micro Average Precision      Recall      F1      Over All Labels
0.531      0.494      0.498
```

For filtered tokens:

```
print("\n Unigram Unfiltered : ")
crossval.cross_validation_PRF(5,filtered_combo_features,labels)
```

```
Unigram Unfiltered :
Each fold size: 31200
Fold 0
      Precision      Recall      F1
0      0.102      0.317      0.154
1      0.162      0.354      0.222
2      0.837      0.623      0.714
3      0.359      0.387      0.373
4      0.144      0.351      0.204
Fold 1
      Precision      Recall      F1
0      0.101      0.304      0.151
1      0.158      0.350      0.218
2      0.845      0.625      0.718
3      0.369      0.405      0.386
4      0.138      0.317      0.193
Fold 2
      Precision      Recall      F1
0      0.100      0.303      0.151
1      0.168      0.348      0.226
2      0.830      0.616      0.707
3      0.345      0.391      0.366
4      0.143      0.321      0.198
Fold 3
      Precision      Recall      F1
0      0.113      0.326      0.167
1      0.164      0.361      0.226
2      0.838      0.614      0.709
3      0.335      0.388      0.360
4      0.147      0.333      0.204
Fold 4
      Precision      Recall      F1
0      0.124      0.316      0.178
1      0.161      0.354      0.221
2      0.839      0.615      0.710
3      0.347      0.396      0.370
4      0.157      0.382      0.222
```

```
Average Accuracy : 0.10851923076923078

Average Precision      Recall      F1      Per Label
0      0.108      0.313      0.160
1      0.162      0.353      0.223
2      0.838      0.618      0.712
3      0.351      0.393      0.371
4      0.146      0.341      0.204

Macro Average Precision Recall      F1      Over All Labels
0.321      0.404      0.334

Label Counts {0: 7069, 1: 27263, 2: 79557, 3: 32912, 4: 9199}
Micro Average Precision Recall      F1      Over All Labels
0.543      0.495      0.499
```

6 Running Naïve Bayes Classifier on Feature sets:

For each of the feature sets generated. Naïve Bayes classifier was run without cross validation.

The highest accuracy for unfiltered tokens was obtained by both SL and SL +LIWC feature sets

The highest accuracy for filtered tokens was obtained by Not feature set

Feature set	Unigram	Bigram	POS	NOT	SL	LIWC	SL +LIWC
Filtered	0.531	0.526	0.520	0.538	0.546	0.540	0.546
Unfiltered	0.513	0.515	0.501	0.560	0.531	0.530	0.531

```
def naivebayesaccuracy(features):
    train_set,test_set = features[int(0.1*len(features)):],
    features[:int(0.1*len(features))]
    classifier = nltk.NaiveBayesClassifier.train(train_set)
    print("\nAccuracy : ")
    print(nltk.classify.accuracy(classifier,test_set),"\n")
    l1 = []
    t1=[]
    for (features,label) in test_set:
        l1.append(label)
        t1.append(classifier.classify(features))
    print(ConfusionMatrix(l1,t1))
```

6.1 Unigram

```
print("\n Unigram filtered : ")
naivebayesaccuracy(filtered_unigram_features)
print("\n Unigram Unfiltered : ")
naivebayesaccuracy(unfiltered_unigram_features)
```

```
Unigram filtered :
Accuracy :
0.5312179487179487

|  0  1  2  3  4 |
+-----+
0 | <27> 107 496 55 7 |
1 | 31 <228>2245 185 24 |
2 | 26 184<7434> 327 24 |
3 | 8 134 2550 <532> 64 |
4 | 4 41 591 210 <66> |
+-----+
(row = reference; col = test)

Unigram Unfiltered :
Accuracy :
0.5135897435897436

|  0  1  2  3  4 |
+-----+
0 | <142> 126 278 94 29 |
1 | 228 <506>1634 309 65 |
2 | 165 518<6725> 464 68 |
3 | 206 275 2119 <518> 165 |
4 | 82 56 502 205 <121> |
+-----+
(row = reference; col = test)
```


6.2 Bigram

```
print("\n Bigram filtered : ")
naivebayesaccuracy(filtered_bigram_features)
print("\n Bigram unfiltered : ")
naivebayesaccuracy(unfiltered_bigram_features)
```

Bigram filtered :						Bigram Unfiltered :					
Accuracy : 0.5268589743589743						Accuracy : 0.5154487179487179					
	0	1	2	3	4		0	1	2	3	4
0	<35>	94	528	76	13	0	<134>	143	338	98	33
1	43	<224>	2225	197	28	1	206	<525>	1618	296	72
2	18	217	<7384>	342	21	2	188	511	<6726>	484	73
3	18	136	2541	<518>	53	3	214	252	2107	<543>	150
4	3	39	575	214	<58>	4	80	53	466	177	<113>

(row = reference; col = test)

6.3 POS:

```
print("\n Pos filtered : ")
naivebayesaccuracy(filtered_pos_features)
print("\n Pos Unfiltered : ")
naivebayesaccuracy(unfiltered_pos_features)
```

Pos filtered :						Pos Unfiltered :					
Accuracy : 0.5201282051282051						Accuracy : 0.5017948717948718					
	0	1	2	3	4		0	1	2	3	4
0	<85>	159	326	85	37	0	<192>	113	238	92	34
1	162	<456>	1729	282	84	1	394	<521>	1430	312	85
2	115	451	<6777>	547	105	2	309	549	<6464>	493	125
3	102	300	2026	<665>	195	3	348	295	1921	<500>	219
4	42	79	412	248	<131>	4	123	70	427	195	<151>

(row = reference; col = test)

1.3 NOT

```
print("\n NOT filtered : ")
naivebayesaccuracy(filtered_not_features)
print("\n NOT unfiltered : ")
naivebayesaccuracy(unfiltered_not_features)
```

NOT filtered :						NOT Unfiltered :					
Accuracy :						Accuracy :					
0.5382692307692307						0.5601282051282052					
	0	1	2	3	4		0	1	2	3	4
0	<387>	227	18	12	9	0	<364>	223	50	8	8
1	621	<1383>	487	205	68	1	679	<1183>	702	137	63
2	416	1435	<4396>	1458	309	2	444	1008	<5223>	1072	267
3	104	248	530	<1701>	674	3	181	178	775	<1446>	677
4	13	18	32	319	<530>	4	20	13	61	296	<522>
(row = reference; col = test)						(row = reference; col = test)					

6.4 SL:

```
print("\n SL filtered : ")
naivebayesaccuracy(filtered_sl_features)
print("\n SL Unfiltered : ")
naivebayesaccuracy(unfiltered_sl_features)
```

SL filtered :						SL Unfiltered :					
Accuracy :						Accuracy :					
0.5467307692307692						0.5316025641025641					
	0	1	2	3	4		0	1	2	3	4
0	<77>	146	363	90	16	0	<165>	158	252	67	27
1	85	<401>	1801	376	50	1	284	<537>	1499	339	83
2	52	336	<6798>	744	65	2	195	500	<6537>	604	104
3	32	174	1819	<1111>	152	3	189	231	1747	<829>	287
4	11	34	302	423	<142>	4	56	42	338	305	<225>
(row = reference; col = test)						(row = reference; col = test)					

1.4 LIWC:

```
print("\n LIWC filtered : ")
naivebayesaccuracy(filtered_liwc_features)
print("\n LIWC Unfiltered : ")
naivebayesaccuracy(unfiltered_liwc_features)
```

```
LIWC filtered :
Accuracy :
0.5407692307692308

| 0 1 2 3 4 |
+-----+
0 | <65> 126 403 86 12 |
1 | 82 <294>1917 375 45 |
2 | 39 250<7018> 636 52 |
3 | 34 128 2067 <938> 121 |
4 | 7 31 394 359 <121> |
+-----+
(row = reference; col = test)

LIWC Unfiltered :
Accuracy :
0.53

| 0 1 2 3 4 |
+-----+
0 | <161> 160 248 74 26 |
1 | 264 <518>1547 327 86 |
2 | 182 495<6604> 562 97 |
3 | 180 253 1810 <789> 251 |
4 | 56 38 395 281 <196> |
+-----+
(row = reference; col = test)
```

6.5 SL+LIWC:

```
print("\n Combined SL LIWC filtered : ")
naivebayesaccuracy(filtered_combo_features)
print("\n Combined SL LIWC Unfiltered : ")
naivebayesaccuracy(unfiltered_combo_features)
```

```
Combined SL LIWC filtered :
Accuracy :
0.546025641025641

| 0 1 2 3 4 |
+-----+
0 | <76> 162 347 97 10 |
1 | 76 <454>1729 408 46 |
2 | 49 382<6700> 808 56 |
3 | 31 204 1748<1155> 150 |
4 | 6 39 285 449 <133> |
+-----+
(row = reference; col = test)

Combined SL LIWC Unfiltered :
Accuracy :
0.5312179487179487

| 0 1 2 3 4 |
+-----+
0 | <172> 156 249 67 25 |
1 | 284 <558>1480 334 86 |
2 | 204 492<6500> 645 99 |
3 | 196 241 1738 <843> 265 |
4 | 53 39 345 315 <214> |
+-----+
(row = reference; col = test)
```

7 Running Random Forest Classifier on all feature sets :

One of the advanced experiments that was chosen was to run the feature sets by saving them into csv files and using these features on a random forest classifier to compare it with the naïve bayes classifier. Taking the internal 90:10 train-test set split, we trained out random forest classifier to calculate the accuracy. For not or negation features, the classifier did not work and gave us memory error. This classifier surprisingly worked very well for unfiltered feature sets as compared to Naïve Bayes. The accuracies were high for filtered data compared to naïve bayes results

Feature set	Unigram	Bigram	POS	SL	LIWC	SL +LIWC
Filtered	0.547	0.546	0.565	0.582	0.567	0.579
Unfiltered	0.567	0.571	0.580	0.605	0.591	0.595

```
def rf(featuresets):
    n = 0.1
    cutoff = int(n*len(featuresets))
    train_set, test_set = featuresets[cutoff:], featuresets[:cutoff]
    classifier_rf = SklearnClassifier(RandomForestClassifier())
    classifier_rf.train(train_set)
    print("Classifier-RandomForest \n")
    print("Accuracy : ",nltk.classify.accuracy(classifier_rf, test_set))
```

For filtered feature sets :

```
Unigram filtered :
Classifier-RandomForest
Accuracy : 0.5472435897435898

Bigram filtered :
Classifier-RandomForest
Accuracy : 0.5467307692307692

Pos filtered :
Classifier-RandomForest
Accuracy : 0.5655128205128205

SL filtered :
Classifier-RandomForest
Accuracy : 0.5820512820512821

LIWC filtered :
Classifier-RandomForest
Accuracy : 0.5673717948717949

Combined SL LIWC filtered :
Classifier-RandomForest
Accuracy : 0.5796794871794871
```

For unfiltered feature sets :

```
Unigram Unfiltered :  
Classifier-RandomForest  
  
Accuracy : 0.5679487179487179  
  
Bigram Unfiltered :  
Classifier-RandomForest  
  
Accuracy : 0.5712820512820512  
  
Pos Unfiltered :  
Classifier-RandomForest  
  
Accuracy : 0.5805769230769231  
  
SL Unfiltered :  
Classifier-RandomForest  
  
Accuracy : 0.6053846153846154  
  
LIWC Unfiltered :  
Classifier-RandomForest  
  
Accuracy : 0.5918589743589744  
  
Combined SL LIWC unfiltered :  
Classifier-RandomForest  
  
Accuracy : 0.5955151541545454
```

Summary/Comparisons:

LIWC feature sets which were not implemented in the course , were implemented and tested out on the reviews dataset.

Combined LIWC and SL features were generated

Random Forest Classifier and Naïve Bayes Classifiers were run, and accuracies were calculated. Random Forest classifier performed better in all cases, but for unfiltered data the accuracy scores were much higher than the filtered featuresets.

Challenges

Although a lot was learnt through trial and error, we faced issues while running the programs as these classifiers are computationally intensive. It was time consuming to run.